

Jordan Ellis

Machine Learning Engineer

jordan.ellis.ml@gmail.com | (403) 555-0147 | Calgary, AB | linkedin.com/in/jordanellisml | github.com/jellis-ml

Summary

Machine learning engineer with 8 years of experience building and shipping production ML systems, with the last 4 focused on applied AI: LLM-based products, retrieval-augmented generation, and model serving at scale. Comfortable owning a system from data pipeline through deployed inference endpoint. Has led a team of 3 engineers and enjoys working close to the product.

Skills

- **Languages:** Python, SQL, Go (basic)
- **ML/AI:** PyTorch, Hugging Face Transformers, scikit-learn, LangChain, LlamaIndex, vector search (FAISS, pgvector, Pinecone), fine-tuning (LoRA/QLoRA), prompt engineering, evaluation harnesses
- **Infra/MLOps:** Docker, Kubernetes, AWS (SageMaker, Bedrock, EC2/S3), Airflow, MLflow, Weights & Biases, CI/CD (GitHub Actions)
- **Data:** Spark, Pandas, dbt, Postgres, Snowflake

Experience

Senior Machine Learning Engineer — Northwind Analytics, Calgary, AB

Mar 2022 – Present

- Built and deployed a retrieval-augmented question-answering system for internal support docs, cutting average ticket resolution time by 34%.
- Led fine-tuning of an open-weight LLM (Llama 3 8B, LoRA) on domain-specific support transcripts, improving answer accuracy over base model by 22% on an internal eval set.
- Designed the team's model evaluation harness (automated + human-in-the-loop scoring), which became the standard gate for all new model releases.
- Mentored 2 junior engineers and led a team of 3 on the applied AI pod; ran sprint planning and technical design reviews.
- Migrated batch inference pipeline from ad-hoc cron jobs to an Airflow + SageMaker setup, reducing pipeline failures by ~90%.

Machine Learning Engineer — Cascade Health Data, Calgary, AB

Jul 2019 – Feb 2022

- Built a patient no-show prediction model (gradient boosted trees) used across 12 clinics, reducing missed appointments by 15% via targeted reminders.
- Owned the feature store and model serving layer (FastAPI + Redis) supporting 4 production ML models with sub-100ms p99 latency.
- Introduced MLflow for experiment tracking and model versioning, replacing spreadsheet-based tracking across the data science team.

- Partnered directly with clinical operations stakeholders to translate business requirements into model objectives and success metrics.

Data Scientist — Prairie Metrics Co., Edmonton, AB

Jun 2017 – Jun 2019

- Built churn prediction and customer segmentation models for a mid-size SaaS client base, informing retention campaigns.
- Automated weekly reporting pipelines in Python/SQL, saving the analytics team ~10 hours per week.
- First engineer to introduce a formal train/test/eval workflow to the team, replacing one-off notebook experiments.

Projects

- **OpenBrief** (personal project): A small open-source tool that generates interview prep briefs from a job posting and resume using a multi-agent LLM pipeline (LangChain + Tavily search). ~200 GitHub stars.
- **eval-kit**: Lightweight Python library for scoring LLM outputs against rubrics, used internally at Northwind and released as open source.

Education

B.Sc. Computer Science — University of Alberta, 2016

Relevant coursework: Machine Learning, Statistics, Distributed Systems

Certifications

- AWS Certified Machine Learning – Specialty (2023)
- DeepLearning.AI: Generative AI with Large Language Models (2023)